

The Collapse of Heterogeneity in Silicon Philosophers

Yuanming Shi, Andreas Haupt
Adobe Inc., Stanford University

Question

Can an LLM stand in for a population of minds?

Setup

- 277 PhilPeople philosophers
- 100 PhilPapers questions
- 7 profile-conditioned LLMs

Test


Does the model induce the same distribution over agent-types as the humans it represents?

A **social world model**— a generative model meant to stand in for a population of human minds— is adequate only if the distribution it induces over **agent-types** matches the population it represents.



Take a picture
for the paper and
data explorer.

Human vs. Claude



Each cell is one philosopher on one question. Humans disagree inside each column. Claude fills columns with one answer.

Paper Figure 3a. No new chart: this is the existing human vs. Sonnet 4.5 plot.